




ENSTA PARIS

# Ray Tracing

LACERDA Filipe  
GUIMARAES Eduardo

# GENERAL PROBLEM

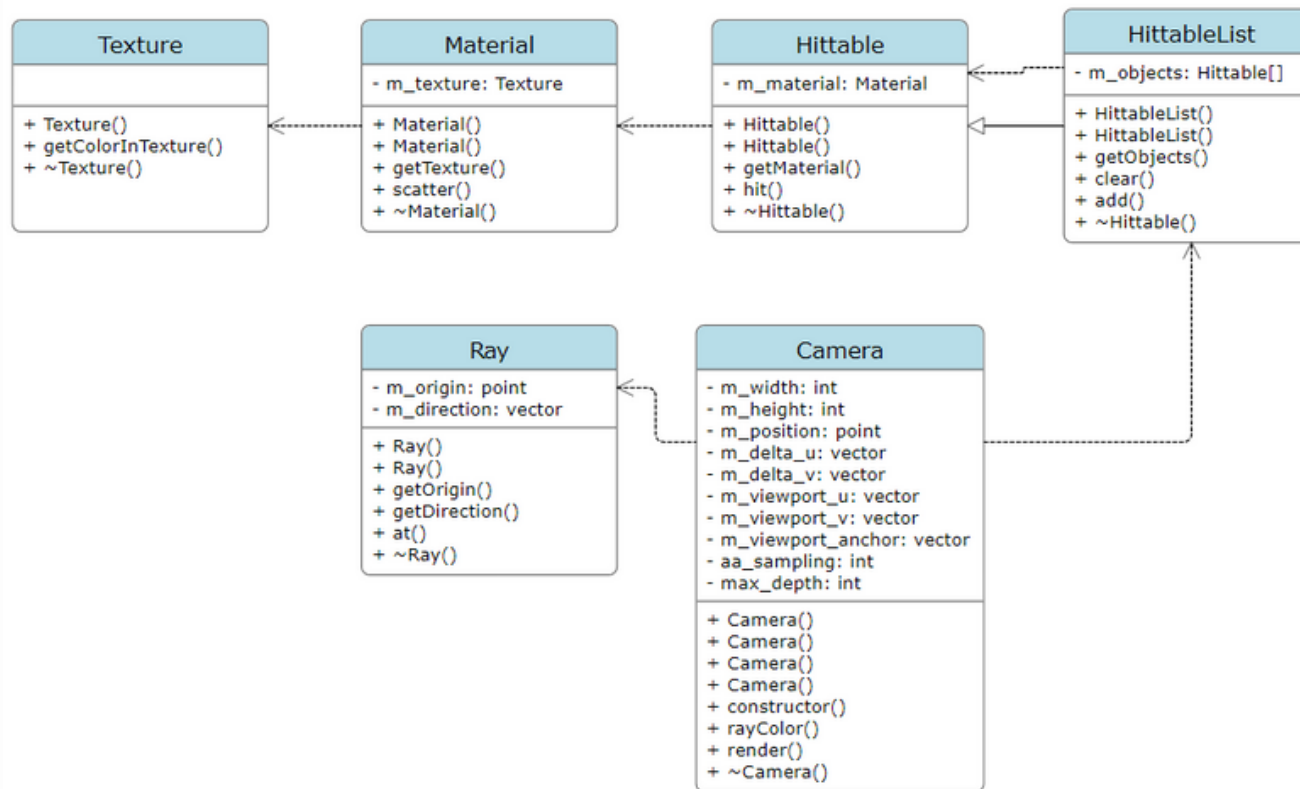
We can draw the base of our graphic motor as:



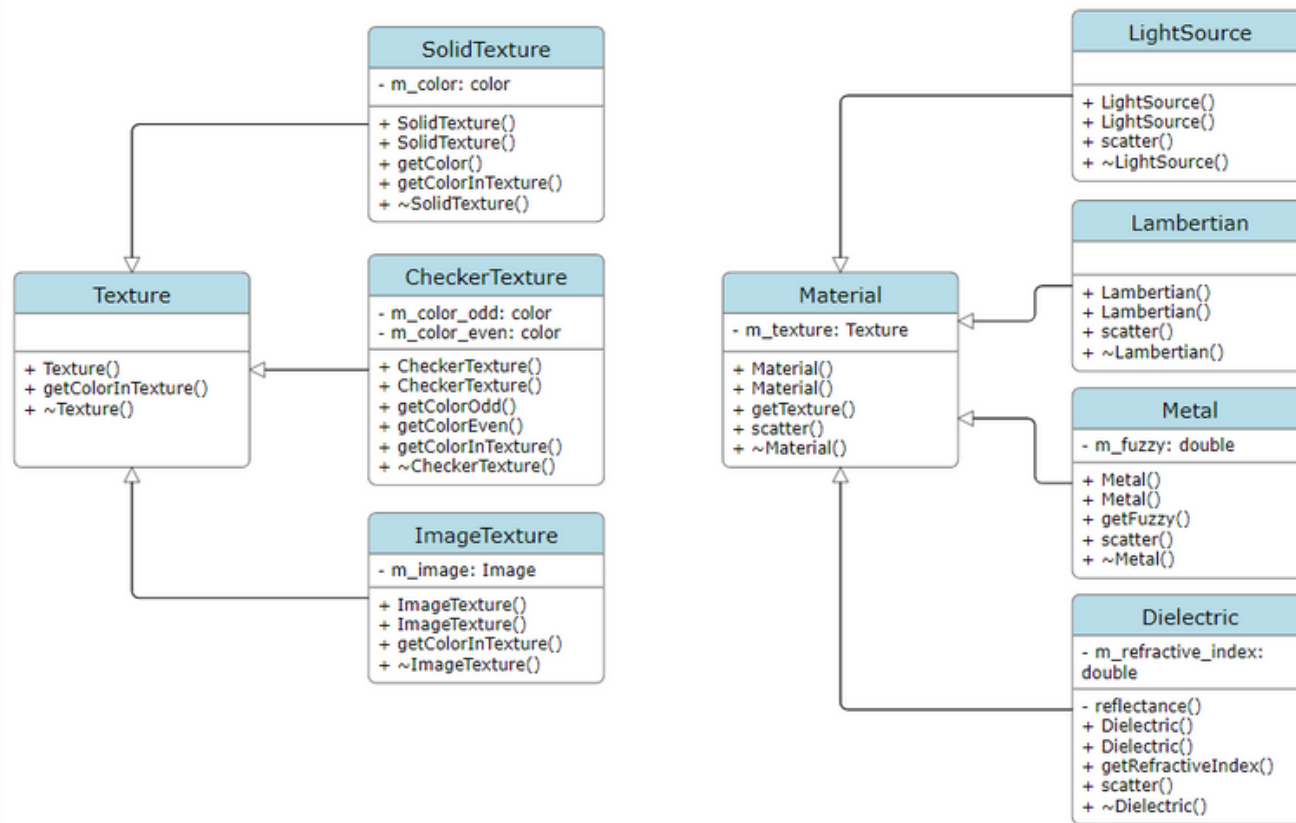
```
<?xml version="1.0" encoding="utf-8"?>

<scene>
  <object geometry="sphere">
    <center>
      <x>1.0</x>
      <y>-1.0</y>
      <z>1.2</z>
    </center>
    <radius>0.38</radius>
    <material appearance="light" texture="solid">
      <albedo>
        <r>1.0</r>
        <g>1.0</g>
        <b>1.0</b>
      </albedo>
    </material>
  </object>
</scene>
```

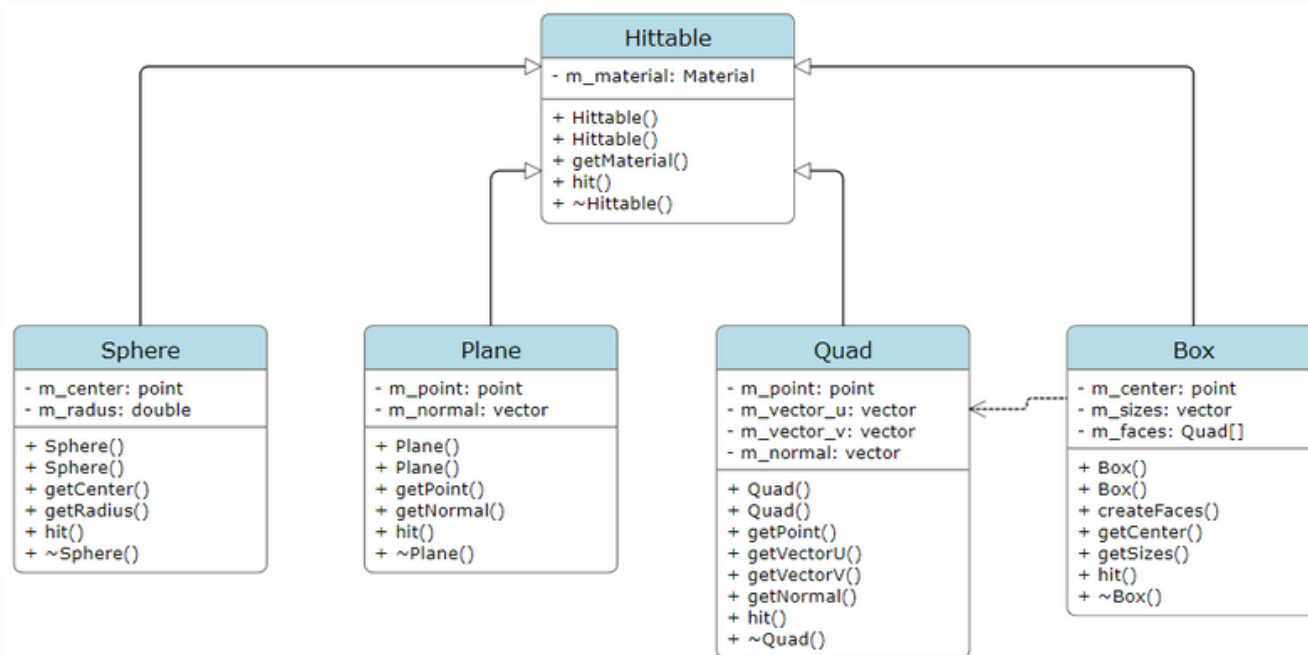
# CLASS RELATIONS




# CLASS RELATIONS



# CLASS RELATIONS







# PERFORMANCE

For certain resolutions and anti-aliasing sampling numbers, our rendering can take several minutes. Where we could start?

- Data and functional parallelization;
- Different approaches to our tracing [2,3], from the statistical approach to anti-aliasing to something like a pre-treatment to ignore some of our rays;
- Code optimization of operations that are repeated multiple times during execution, like square roots and normalization.



# FAST INVERSE ROOT

- Inverse Square Root is an expensive operation, made everytime a vector/ray is normalized;
- A faster interactive version was used for Quake III:

```
inline float q_sqrt(float number){
    long i;
    float x2,y;
    const float threehalfs = 1.5f;

    x2 = number * 0.5f;
    y = number;
    i = * ( long * ) &y; // evil floating point bit level hacking
    i = 0x5f3759df - (i >> 1); // what the ****?
    y = * ( float * ) &i;
    y = y * ( threehalfs - ( x2 * y * y )); // 1st iteration
    y = y * ( threehalfs - ( x2 * y * y )); // 2nd iteration

    return y;
}
```





# PARALLELIZATION

## OPTIONS:


- GPU with CUDA: Fast, but big changes would be made to the code;
- OpenMPI: Task balance is not that trivial;
- OpenMP: We have a simple loop, so task balance can be made by the API's own scheduler.

## OPENMP IT IS!

```
num_threads = omp_get_num_procs();//gets total number of threads
omp_set_dynamic(0);//lets code change number of max threds used in parallel block
omp_set_num_threads(num_threads);//sets number of threads used in a parallel block
#pragma omp parallel for private(i) schedule(dynamic,10)
for (j = 0; j < m_height; j++) {
    for ( i = 0; i < m_width; i++) {
```

# PARALLELIZATION

- Images got at least 2x faster(the optimum parameters vary from sytem to system;
- Easy to mix it with MPI, in the case we have multiple multiprocessors in a single system.





# PERFORMANCE

WHAT COULD BE DONE TO MAKE IT FASTER?

- The parallelization with CPU is nice, but using a GPU could make things even 10x faster [3]. We could also use an approach that mixes both [4];
- Even though we made the inverser square root faster, other inefficient operations remain (normal square root and antialiasing sampling for an example);
- Some smarter aproaches to raytracing could also be made, inspired by what they do in videogames, like leaving ray tracing to shadows and reflections and using rasterization for the rest;

# References

- [1] J. Q. Michael, "Parallel Programming in C with MPI and OpenMP" MC GRAW HILL, 2004.
- [2] S. Peter, D. B. Trevor and H. Steve, "Ray Tracing in One Weekend Series", 2018-2023, free access in <https://raytracing.github.io/>.
- [3] A. Roger, "Accelerated Ray Tracing in One Weekend in CUDA", 2018, free access in <https://developer.nvidia.com/blog/accelerated-ray-tracing-cuda/>.
- [4] M. Evgeny, "Practical Tips for Optimizing Ray Tracing", 2023, free access in <https://developer.nvidia.com/blog/practical-tips-for-optimizing-ray-tracing/>.